

Ns Tutorial: Case Studies

John Heidemann (USC/ISI)

Polly Huang (ETH Zurich)

March 14, 2002


Road Map

- Simple examples ← **Provide an entry point**
 - TCP
 - web traffic
- Case Study ← **Show case ns's functionality and relevance to the CN program**
 - Impact of HTTP and TCP parameters to Web performance
 - Hidden structure behind aggregated Web traffic

Presentation Style

- Slides
- Script walk-through
- Live demos with nam (Network AniMator)

Example I: TCP



```
set ns [new Simulator]
```

```
set n0 [$ns node]
```

```
set n1 [$ns node]
```

```
$ns duplex-link $n0 $n1 1.5Mb 10ms  
    DropTail
```

```
set tcp [new Agent/TCP]
```

```
set tcpsink [new Agent/TCPSink]
```

```
$ns attach-agent $n0 $tcp
```

```
$ns attach-agent $n1 $tcpsink
```

```
$ns connect $tcp $tcpsink
```

```
set ftp [new Application/FTP]
```

```
$ftp attach-agent $tcp
```

```
$ns at 0.2 "$ftp start"
```

```
$ns at 1.2 "exit"
```

```
$ns run
```

Example II: Web Traffic

- A Web session – a series of page downloads
 - Number of pages
 - Inter-page time
 - Page size (number of embedded objects)
 - Inter-object time
 - Object size (KB)
- 5 random variables

Case Study I: Web Performance

- Impact of TCP and HTTP parameters
- Try to answer:
 - Will the proposed changes work in a variety of conditions?
 - Should TCP Sack be deployed?
 - Should persistency or pipelining be deployed?
 - Which parameters are more cost effective to tune?

Methodology

- Methodology
 - Select performance critical parameters
 - Use most commonly used values as the base case
 - Tune parameter values to compare to the base case
- Toward a systematic and exhaustive evaluation
- Enabled by ns
 - rich library of workload and protocol implementations
 - Contributed code from a huge user/developer community

Parameters and Values


TCP

- Packet size
 - 576, 1460
- Delayed ack
 - on, off
- Congestion avoidance
 - NewReno, Tahoe, Reno, Sack
- Initial retransmission timeout
 - 3, 6 sec
- Timer granularity
 - 100, 500 msec
- Timestamp option
 - on, off
- Initial window size
 - 2, 4

HTTP


- Connection type
 - persistent, simple, pipelined
- Number of parallel connections
 - 2, 1, 4

Page Download Time – TCP



**Sack, NewReno, Reno, Tahoe, gradually better
Timer-related parameters, no significant impact**

Page Download Time - HTTP



**Simple, persistent, and pipelined connections,
gradually better**

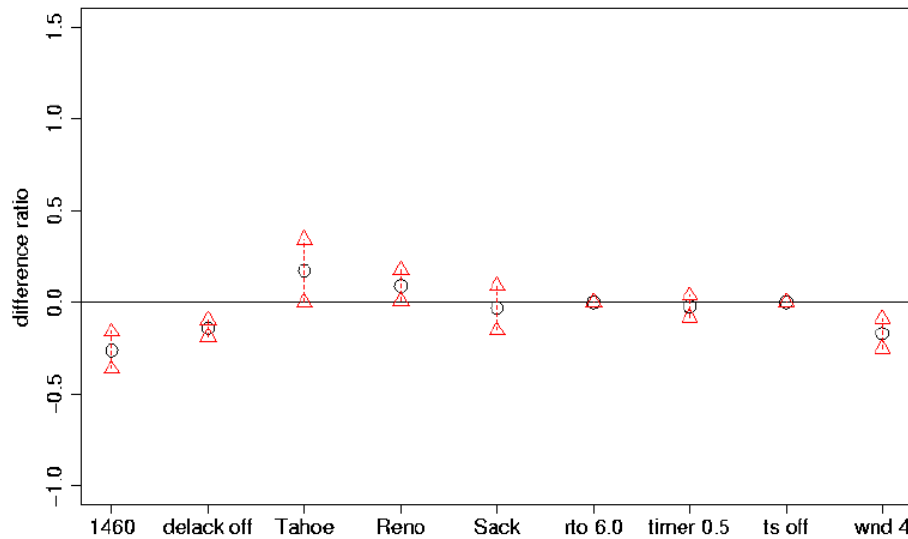
**Higher the number of parallel connections,
Smaller the range of improvement**

TCP vs. HTTP


TCP

HTTP

TCP-- low-loss case: page-download.summary col= 4/4



HTTP-- low-loss case: page-download.summary col= 4/4



Impact of TCP Sack is relatively small.



Low vs. High loss - TCP

Low loss

High loss

TCP--- low-loss case: page-download.summary col= 4/4


TCP--- high-loss case: page-download.summary col= 4/4



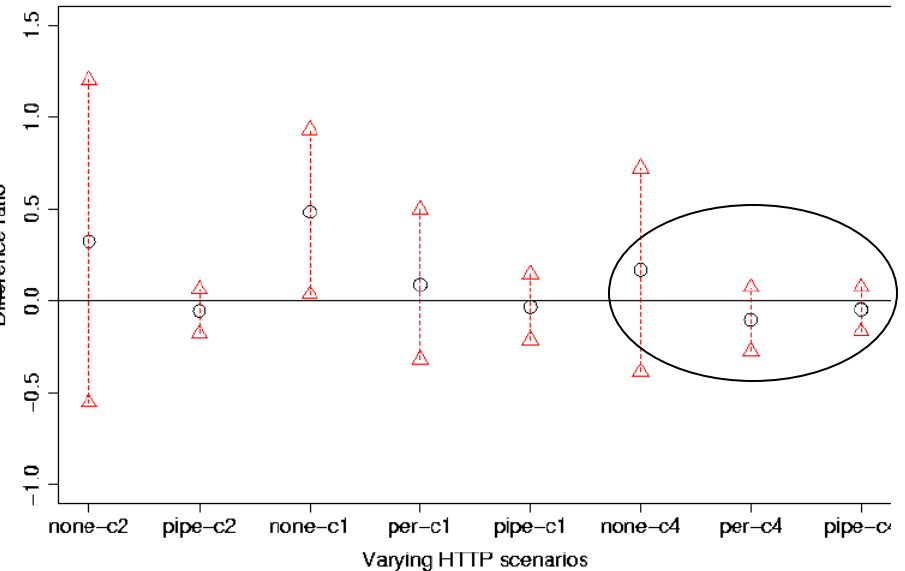
That tiny bit of advantage in TCP Sack disappears in high-loss case.

Low vs. High loss - HTTP

HTTP-- low-loss case: page-download.summary col= 4/4



HTTP-- high-loss case: page-download.summary col= 4/4



**Pipelining loses its advantage when
of parallel connections is high.**

Preliminary Findings

- Will the proposed changes work in a variety of conditions?
 - Not really
 - TCP Sack and HTTP pipelining
- Should TCP Sack be deployed?
 - Maybe not, if deployment cost is high
- Should persistency or pipelining be deployed?
 - Maybe yes, but doesn't make sense to work with too many parallel connections

The Real Message











- Design decisions need to be validated in the context of the Internet.
- Layers of protocols, tremendous amount of unknown dynamics
- Simulation tools like ns can help us track the complexity (within a layer or across layers)
- Ns en-powers such studies
 - A rich library base
 - A large community contributing to the base

Case Study II: Web Traffic

- Web traffic is not exact self-similar
- How does it diverge from exact self-similar?
- Why is there this divergence?


Self-similarity

- Distributions of #packets/time unit look alike in different time scale

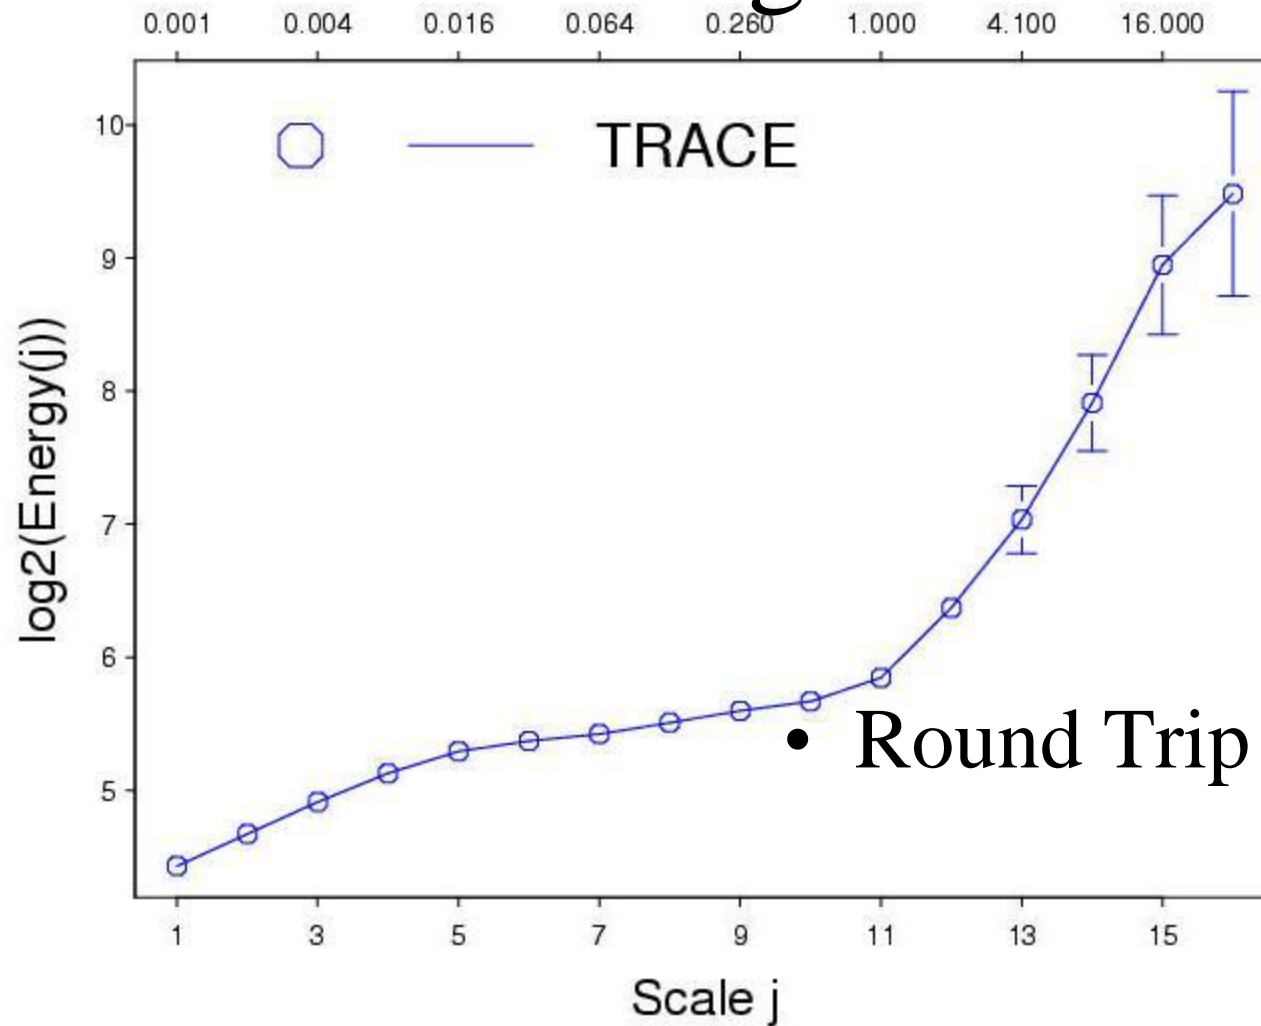


Wavelet Analysis

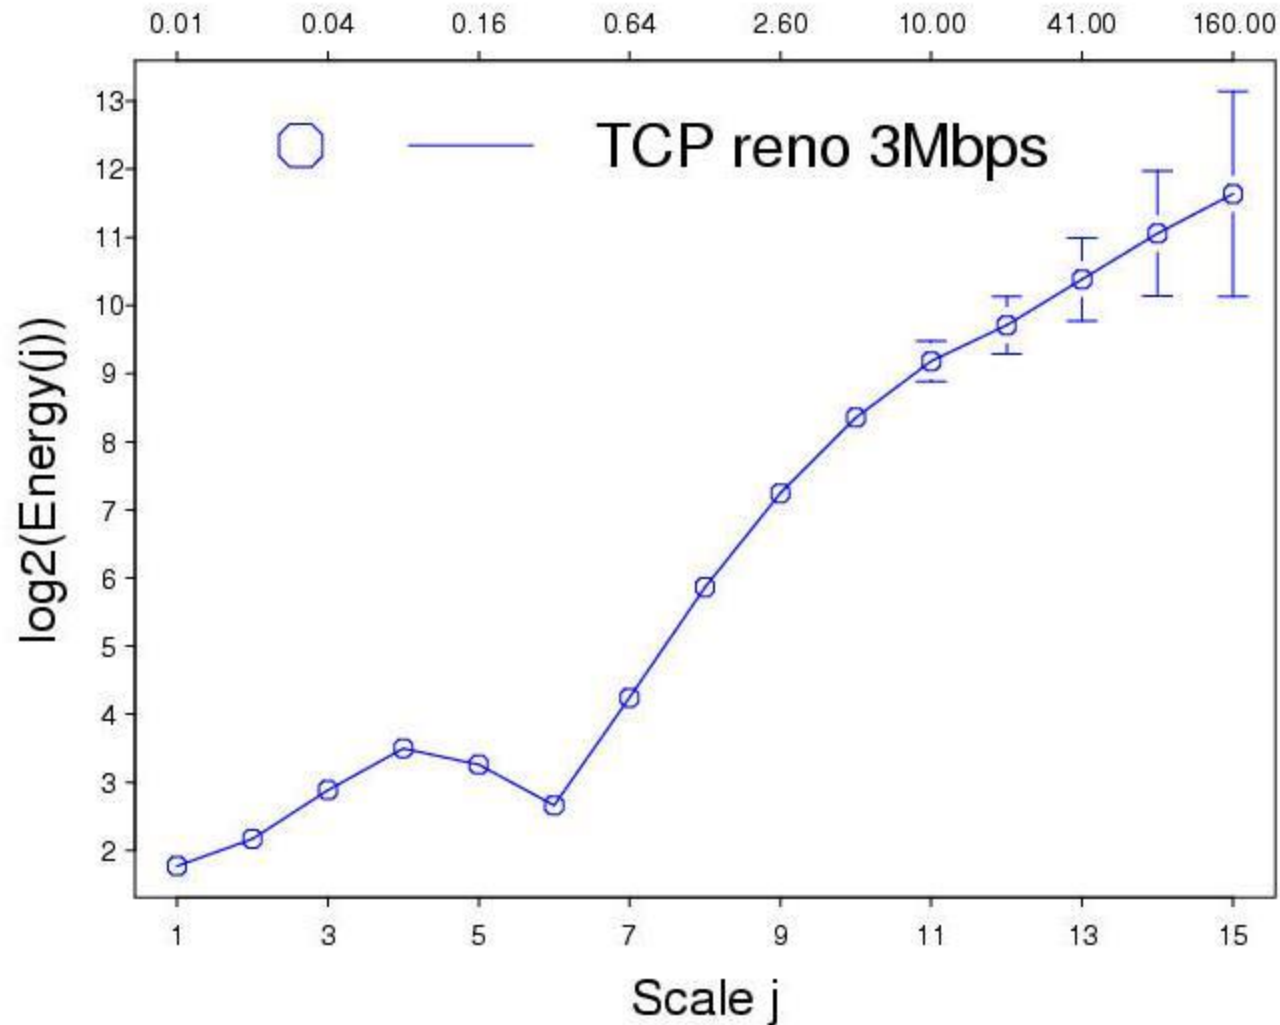
- FFT - frequency decomposition d_j
- WT - frequency and time decomposition $d_{j,k}$
- $\sum_k (d_{j,k}^2) / N_j \equiv E_j = 2^{j(2H-1)} C$
- $\log_2 E_j = (2H-1) j + \log_2 C$



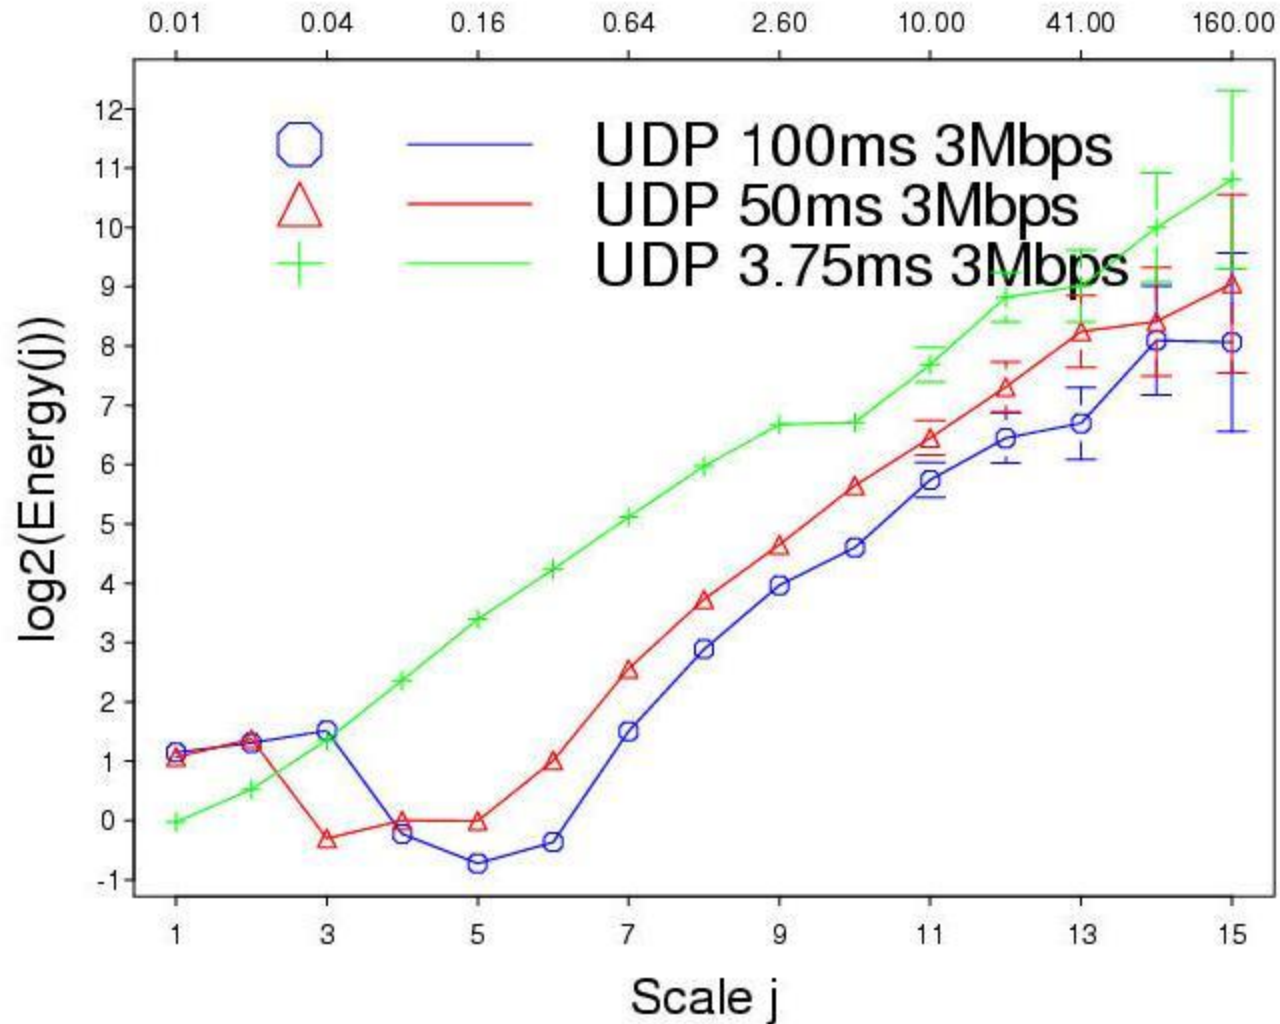
Global Scaling - Trace



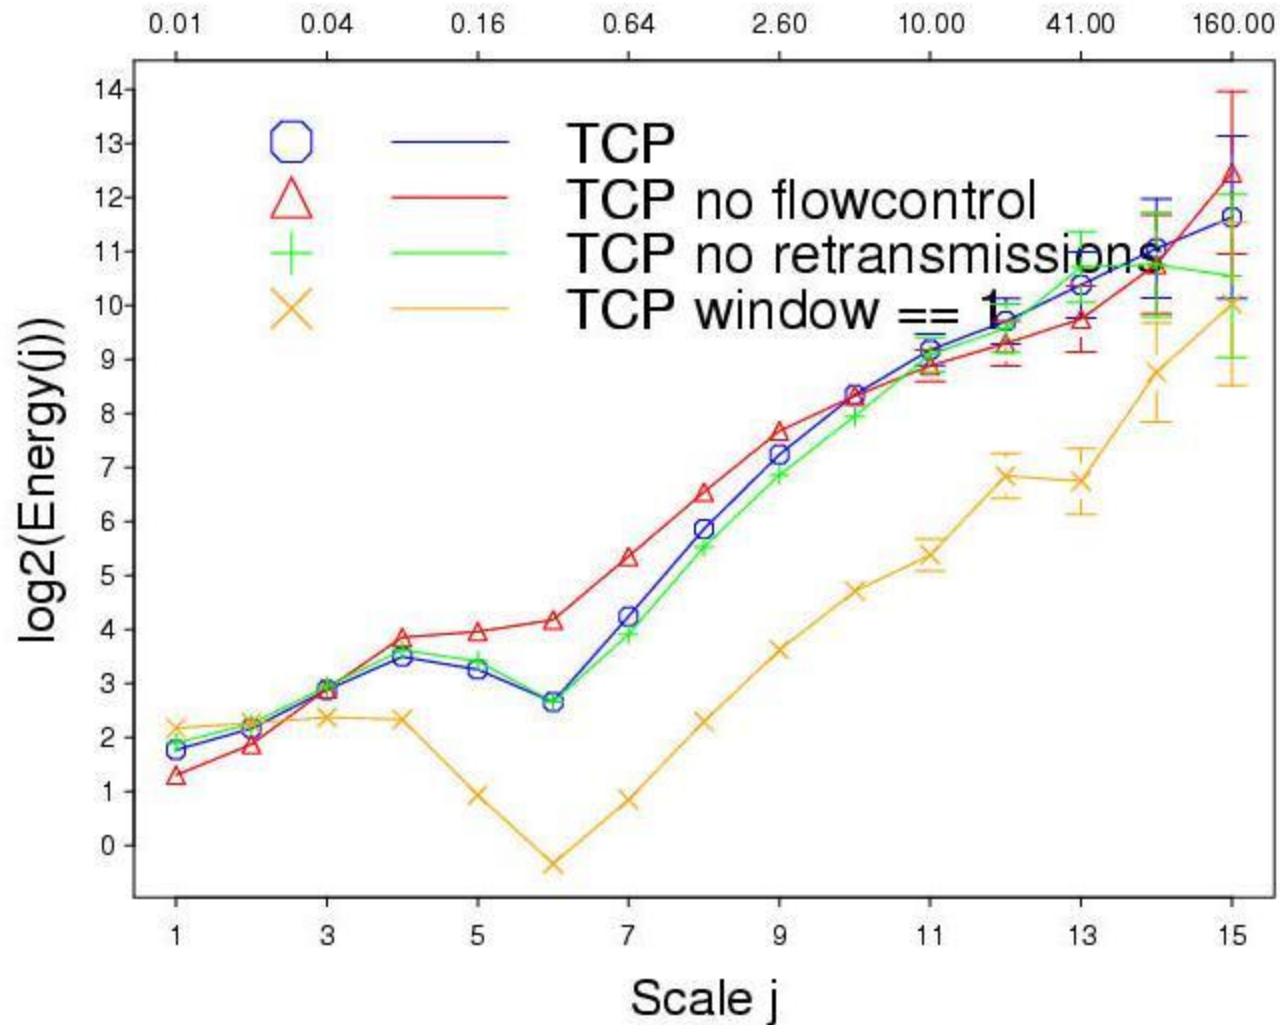
Global Scaling - Simulation



UDP



TCP



Findings

- Periodicity emerges at round-trip time scales
- That periodicity dominates the traffic behavior at those scales
- TCP ack clocking plays a critical role
- Need to be cautious when to use or not use mathematical self-similar models

The Real Message

- Proposed (traffic) models need to be validated in the context of the Internet.
- Mechanisms can influence Internet characteristics in a surprising way
- Simulation tools like ns can help us track the implicit complexity
- Ns en-powers such studies
 - A rich library base
 - A large community contributing to the base

Concluding remarks

- Learning ns
 - video recording (huang@tik.ee.ethz.ch)
 - on-line tutorials (audio and slides)
 - tons of info from the ns web site
- Research with ns
 - promote sharing and confidence